
Big Data Platforms - Recitation Workshops

Workshop II
Data & Cloud Storage

21 December 2022

Yarden Fogel & Roei Arpaly



Data & Cloud Storage Agenda

Data & Engines


- Data Lakes, Data Warehouse
- SQL, NoSQL and Databases
- Files (text, csv, image, json, yaml, etc)
- Parquets (partitions, predicate pushdown, columnar storage)
- Big data engines, query optimisation, scaleup
- PySpark

Storage

- Cloud Storage, File System and Object Storage
- Storing Models (Pickle, Joblib)
- Model Store / Feature Store


Data Lake and Data Warehouse

- A **data lake** is a repository of data stored in its raw format.
- A data lake can include:
 - **structured** data from relational databases (rows and columns)
 - **semi-structured** data (logs, CSV, XML, JSON)
 - **unstructured** data (images, audio, video, binary data)
- Most of the data lakes nowadays are resides in a cloud storage (using the service of vendors such as Amazon, Microsoft or Google).



Data Lake and Data Warehouse (Cont'd)

- A **data warehouse** (DW or DWH) is a system used for reporting and data analysis and is considered a core component of business intelligence (BI).
- They store processed (transformed) historical data in one single place that is used for tasks such as reporting, visualization, advanced analytics and machine learning.
- ETL and ELT are the two main approaches used to build a DW system (will be discussed later).



Relational and Non-Relational databases

- **SQL** - Structured Query Language, working with relational databases which store data in the form of schema, tables, columns, and rows.
 - Hard to scale, hard to manage data with different data models
 - Examples: SQLite, MySQL, PostgreSQL, MS SQL, etc
- **NoSQL** - a NoSQL database is one that is less structured, allows for more flexibility and adaptability. No predefined relationships or constraints in the database.
 - Popular types: key-value, document, graph and column.
 - MongoDB, DynamoDB, etc

Big Data Engines

- A Big data engine is a processing engine which is responsible for processing data, usually retrieved from cloud storage.
- Big data engines utilize a distributed parallel programming framework that enables processing very large amounts of data distributed across multiple nodes.
- **Master/slave** (or coordinator/worker) relationship
- **Optimisation** (analyzer, execution-plan, optimisation (predicate-pushdown etc))

Big Data Engines (Cont'd)

- Why is sorting an expensive task even with multiple workers (parallelisation)?
- Why is limiting the amount of rows after sorting more efficient in parallel?

Sorting example:

- Master with 10 Workers, 100M rows.

Big Data Engines (Cont'd)

BI tool

MicroStrategy



looker



Superset



+ a b l e a u

Engine

Presto
Cluster



Coordinator

Worker

Worker

Worker

Worker

Worker



Database /

File system /

Object Storage



amazon
REDSHIFT

Microsoft
SQL Server

MySQL

PostgreSQL

accumulo

kafka

elastic

mongoDB

pinot

redis

druid

cassandra

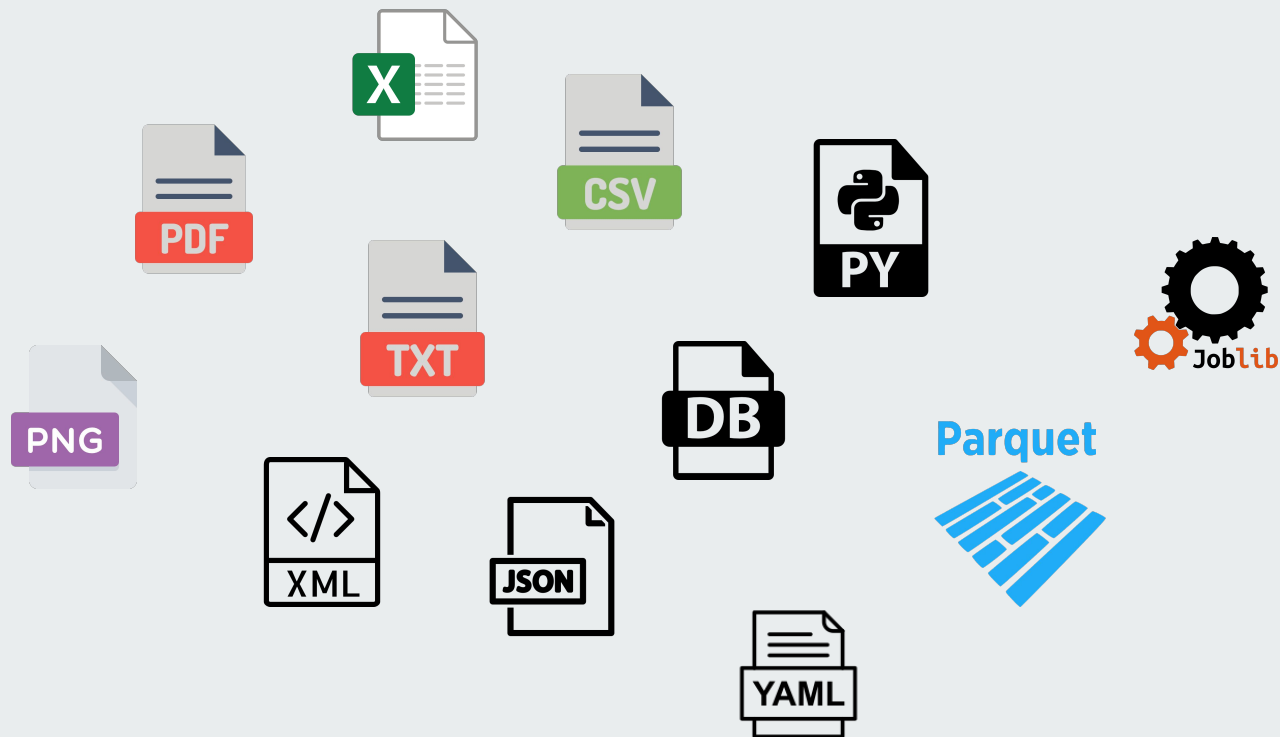
hadoop
HDFS

<https://prestodb.io/>

PySpark

- **PySpark SQL:** Spark SQL is a Spark module for structured data processing. It provides a programming abstraction called DataFrame (distributed datasets) and can also act as distributed SQL query engine.
 - Spark allows working with pandas API to scale your pandas workload out.
- **Pyspark MLlib:** Built on top of Spark, MLlib is a scalable machine learning library that provides a uniform set of high-level APIs that help users create and tune practical machine learning pipelines.

Files



...and many
more...

XML, JSON, YAML

XML

```
<note>
  <to>BDP students</to>
  <from>BDP staff</from>
  <heading>Recitation hours</heading>
  <body>Recitation will start today at 19:15</body>
</note>
```

JSON

```
{
  "to": "BDP students",
  "from": "BDP staff",
  "heading": "Recitation hours",
  "body": "Recitation will start today at 19:15"
}
```

YAML

```
to: BDP students
from: BDP staff
heading: Recitation hours
body: Recitation will start today at 19:15
```

Parquet

- **Apache Parquet** is an open source, column-oriented data file format for retrieval and storage of large amounts of data. It provides efficient data compression and encoding schemes.
- **Partitions:** A partition column is a column that the data files do not store values for. Instead, when writing the files, they will be divided (grouped) into folders / files based on partition column values.
- **Predicate Pushdown:** Predicate pushdown deals with what values will be scanned in the parquet files. If apply a filter on a specific column, to only return records with a specific value, the predicate push down will make parquet read only blocks that may contain this value.
- **Columnar Storage:** Parquet files allow the user to retrieve only the selected columns, reducing the amount of data transferred.

Parquet (Cont'd)

- **Parquet files management:** Parquet files are immutable, therefore only CREATE and OVERWRITE methods are supported.
- **Overwrite partitions:** If a parquet file is corrupted or is missing data, it can be overwritten. Hence, the file name should be easy to identify.
- **Time partitions:** Usually, one of the partition hierarchies should be time-based. Imagine you have a data pipeline which runs every day, in case there was a problem with a specific run, re-run will overwrite the files under the relevant time partition.
 - A dtstring (datetime in string format) column can be created for partitioning.

Query Optimization

- **Using partitions:** query over limited amount of folders of parquet files.
- **Predicate pushdown:** retrieve only limited information
- **Columnar format:** try not to use SELECT *
- **ORDER BY:** sorting only in the final query
- **JOIN:** small table to the bigger table
- **GROUP BY:** group by higher cardinality first, if possible join after the aggregation
- **DISTINCT / UNION:** use only if necessary
- **LIMIT:** retrieve less amount of data, important especially when sorting
- **APPROX aggregate functions:** more efficient functions (on behalf of accuracy)
- **REGEX:** more efficient than the LIKE operator

* Most engines try to automatically detect opportunities to optimize queries based on some of those examples and additional techniques

Pickle, Joblib

- **Pickle:** The pickle module implements binary protocols for serializing and de-serializing a Python object structure. “Pickling” is the process whereby a Python object hierarchy is converted into a byte stream (serialization), and “unpickling” is the inverse operation.
<https://docs.python.org/3/library/pickle.html>
- **Joblib:** Joblib is a set of tools to provide lightweight pipelining in Python.
In particular: transparent disk-caching of functions and lazy re-evaluation (memoize pattern), easy simple parallel computing. Joblib is optimized to be fast and robust on large data in particular and has specific optimizations for numpy arrays.
<https://joblib.readthedocs.io/en/latest/>

Pickle - examples from the docs

Examples

For the simplest code, use the `dump()` and `load()` functions.

```
import pickle

# An arbitrary collection of objects supported by pickle.
data = {
    'a': [1, 2.0, 3+4j],
    'b': ("character string", b"byte string"),
    'c': {None, True, False}
}

with open('data.pickle', 'wb') as f:
    # Pickle the 'data' dictionary using the highest protocol available.
    pickle.dump(data, f, pickle.HIGHEST_PROTOCOL)
```

The following example reads the resulting pickled data.

```
import pickle

with open('data.pickle', 'rb') as f:
    # The protocol version used is detected automatically, so we do not
    # have to specify it.
    data = pickle.load(f)
```


Joblib - example from the docs

A simple example

First create a temporary directory:

```
>>> from tempfile import mkdtemp
>>> savedir = mkdtemp()
>>> import os
>>> filename = os.path.join(savedir, 'test.joblib')
```

Then create an object to be persisted:

```
>>> import numpy as np
>>> to_persist = [('a', [1, 2, 3]), ('b', np.arange(10))]
```

which is saved into *filename*:

```
>>> import joblib
>>> joblib.dump(to_persist, filename)
['...test.joblib']
```

The object can then be reloaded from the file:

```
>>> joblib.load(filename)
[('a', [1, 2, 3]), ('b', array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9]))]
```

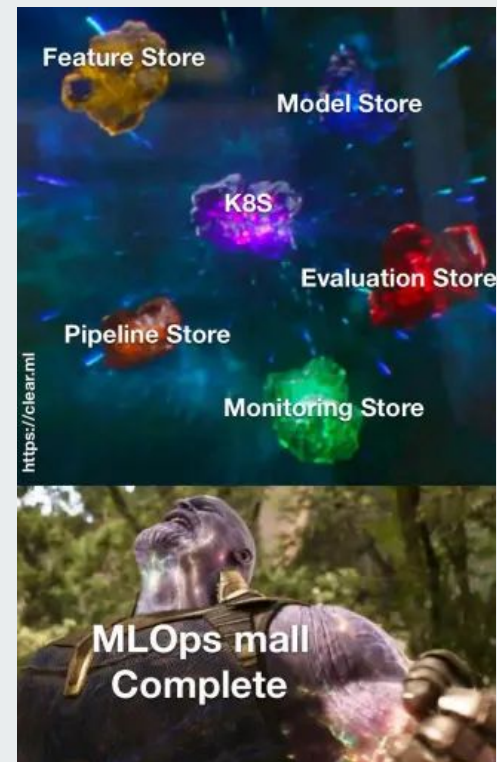
Cloud Storage, File System and Object Storage

- **Cloud Storage:** stored in logical pools, on remotely multiple physical servers, usually owned and managed by a hosting company.
- **File System:** data structure that a operating system uses to control how data is stored and retrieved. The data is separated into pieces (folders and files) with a name for easy identification.
- **Object Storage:** Object storage manages data as objects, as opposed to other storage architectures like file systems which manages data as a file hierarchy. Each object typically includes the data itself, a variable amount of metadata, and a globally unique identifier.

Model Store

- **Model Store:** A model store is a central storage for ml engineers / data scientists to push, pull, version and manage their models and experiments, including the model files, artifacts, and metadata.

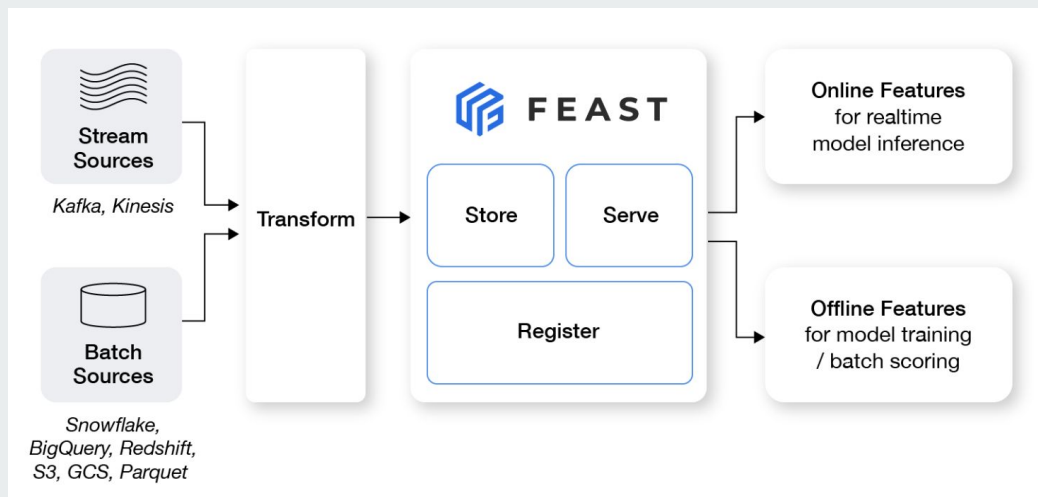
A centralized storage for models allows to compare new models against existing deployed versions, track model performance over time, and more..



Source: <https://imgur.com/gallery/w9DnEA7>

Feature Store

- **Feature Store:** A feature store is an emerging infrastructure for organizing data for Machine Learning use cases, especially for real-time purposes. The features can be shared across different ML pipelines.



Questions?



Open the Notebook